

[Java] Coding Challenge: Candle Aggregation Service

Overview

You are building a backend **Java service** that listens to a continuous stream of **bid/ask market data**, aggregates this data into **candlestick (OHLC)** chart format, and exposes a history API to provide this data to frontend charting libraries such as **Trading View Lightweight Charts** or custom dashboards. The service should support multiple symbols and timeframes and be capable of storing and serving historical candle data efficiently.

Functional Requirements

1. Stream Ingestion:

- Ingest a real-time or simulated stream of market data events:

```
record BidAskEvent(String symbol, double bid, double ask,  
long timestamp) {}
```

- Use a random generator, scheduled tasks, or real data sources like Kafka or WebSocket feeds.

2. Candlestick Aggregation:

- Aggregate bid/ask events into OHLC candles per:

- Symbol (e.g., BTC-USD, ETH-USD)
 - Interval (e.g., 1s, 5s, 1m, 15m, 1h)

- Output structure:

```
record Candle(long time, double open, double high, double  
low, double close, long volume) {}
```

- Volume can be synthetic (number of ticks) or event-derived.

3. Data Storage:

- Store candle data for each `(symbol, interval, timestamp)` tuple.
- Minimum: Use in-memory storage (e.g., ConcurrentMap or H2).
- Preferred: Use a time-series-optimized database (PostgreSQL, TimescaleDB).

4. History API Endpoint:

- Provide historical candle data through the following REST endpoint:

```
GET /history?symbol=BTC-  
USD&interval=1m&from=1620000000&to=1620000600
```

- Expected response:

```
1 {  
2   "s": "ok",  
3   "t": [1620000000, 1620000600, ...],    // time  
4   "o": [29500.5, 29501.0, ...],          // open  
5   "h": [29510.0, 29505.0, ...],          // high  
6   "l": [29490.0, 29500.0, ...],          // low  
7   "c": [29505.0, 29502.0, ...],          // close  
8   "v": [10, 8, ...]                      // volume  
9 }  
10
```

- Timestamps are in UNIX seconds.

Non-Functional Requirements

1. Performance:

- Able to handle high-frequency bid/ask updates (e.g., multiple per second per symbol).
- Minimize latency in real-time candle generation and history responses.

2. Scalability:

- Support multiple symbols and intervals concurrently without blocking.
- Handle time-aligned candle generation even with slight event delays.

3. Reliability:

- Thread-safe aggregation and storage.
- Support safe startup, shutdown, and replaying of missed data if needed.

4. Maintainability:

- Clean and modular code.
- Unit and integration tests for the core logic (especially aggregation).

5. Observability:

- Logging for received events and candle creation.
- Health check endpoints or logging to confirm proper operation.

6. Extensibility:

- Easy to add new timeframes or symbols.
- Candle calculation logic is decoupled from data ingestion source.

Submission Details

Please submit your solution via a **GitHub repository link**. The repository should include:

- A **clear README.md** file with:

- Project overview
- Any assumptions or trade-offs made
- Instructions for running tests
- Notes on bonus features (if implemented)